

Response to peer reviews

September 26, 2021

Editor

Reviews are generally supportive, but bring up several issues that need to be addressed. Once addressed, please submit the revised manuscript with an item-by-item reply/rebuttal.

We thank the editor and reviewers for their comments and the opportunity to make revisions. We have made a number of revisions, described in detail below. We believe these changes have strengthened the manuscript, and describe them in more detail below. We quote below the Reviewer comment and provide our response in this color.

Reviewer 1

1. In their presentation of of evaluation metrics they do not discuss stability. For instance in [7] does the choice of the weights, using τ for $x \geq 0$ rather than some monotone function say some power τ^α make a difference? Since as they indicate the measure is used for fitting as well as ascertainment, some stability investigation seems appropriate. I suspect here is no great effect but...

The reviewer is correct: the choice of τ rather than a monotone function of τ will make a difference. But it seems that we should have been more clear in our description of WIS. The weights used here are those used to evaluate forecasts submitted to the COVID-19 Forecast Hub and the CDC. More broadly, this is a standard metric in the forecasting community (just as Mean Squared Forecast Error is standard), rather than an “author decision”. In principle, we would expect that different weights may alter the

conclusions, though likely not substantively. Applying a monotone function to τ only effectively changes the quantile of interest, resulting in a mismatch between coverage you tried to get and the coverage WIS is evaluating. Applying the function to τ and $(1 - \tau)$ results in an asymmetry, penalizing forecasts that miss on one side more than on the other.

- We have added a sentence in the Evaluation Metrics section emphasizing that this metric is standard.

2. They do not discuss the possible effects of vaccination rates and other factors, though this may be one of the areas of localization they propose to investigate.

During the period discussed in the main paper, vaccinations hadn't really taken off. On Dec 31 2000, when the main evaluation period ends, according to CDC data, about 1% of the population had received 1 dose. That said, we should probably be more specific in the paragraph on lines 31-40: our goal is to use very simple models that do not account for extra information (vaccinations but also strains of the virus, "super-spreader" events, etc.).

- We have added a sentence around lines 30–40 emphasizing the choice to ignore such factors.
- We also added some text to the Prediction Period section stating that our end date (31 December) coincides with the introduction of vaccinations.

3. This points to what is perhaps a more major issue. Can they identify in quantitative terms the consequences of the additional improvements in forecasting provided by these methods in terms of gains in resources available to deal with the epidemic? Unfortunately as they state, in the crucial periods of sharply rising cases, the additional benefit seems least.

This is a good point, and we have made a few modifications described in more detail below (Reviewer 2, Point 3).

4. Despite the fact that a detailed description of terms in the text is impossible, it would I think be helpful to say if there is any significance to the names they associate with their first 2 and 4th indicators.

It seems that the Reviewer is referring to CHNG-COVID and DV-CLI. These are described in the Methods Section, Signals and Locations Subsection. To be more

specific, Change Healthcare is a large healthcare insurance claims processor. The difference between “-CLI” and “-COVID” is that the first measures insurance claims related to symptoms consistent with COVID while the second requires a medical diagnosis (through testing or presumed positive by the medical provider). The “DV” modifier corresponds to different claims processors (not Change Healthcare) who wish to remain anonymous.

TODO Do we want to alter the text at all?

5. What are Change Healthcare claims..as opposed to others?

Hopefully the discussion to point 4 above has cleared up any confusion.

Reviewer 2

(items numbered by the editor for convenience of reference)

Statistical or Methodological Comments:

1. Is this really a stationary phenomenon? Over the timerange extending from pre-lockdown to today, it certainly isn't. The method in this paper of presenting model prediction errors makes it impossible to see this. People in econometrics look at cumsums of normalized prediction errors. You can see change points and trends very easily.

There is little reason to believe that this data is stationary. The models however are trained and evaluated over reasonably short periods (2-3 months), over which time the behavior may well be stationary. We have added a section to the appendix that uses the cumulative mean as suggested (the sum is difficult to see due to the scaling). With some exceptions at the beginning of the evaluation period, the relative ordering remains fairly consistent. The AR model is easily the worst.

TODO I have added the code and a caption for the figure to the Supplement, but no corresponding text. It seems a fairly minor addition, not necessarily deserving its own section.

2. Are the results in this paper statistically significant? I see no discussion of this question, which seems unbelievable for me to be saying, given the authorship. And yet here I am.

This is an important point that we should have addressed more fully. Conventional methods for assessing statistical significance that we explored were overconfident, to the point of being unbelievable. To a large degree, this is a “huge n, tiny p” problem so even bootstrap analyses result in extremely narrow intervals. Essentially, “everything is significant”.

We also have strong reasons for not using model-based approaches (one being the nonstationarity the Reviewer mentions above, but also the belief that these models are wrong). We felt that the related investigation in 2B was more pertinent.

To respond more directly,

- we have added a section to the Supplement (and point to it around line 85 in the main text), that performs a Sign Test for differences in accuracy. P-values there tend to cluster near zero suggesting again that these improvements are statistically significant.
- We also use a Diebold-Mariano test for comparing forecast accuracy. The results for this test are more mixed, likely due to nonstationarity the Reviewer mentions above (the indicators are more helpful in some periods than others, see also the upswing/downswing discussion).

3. Are the results in this paper practically significant? I see no discussion of this question, which seems unbelievable for me to be saying, given that the authors refuse to consider statistical significance. And yet here I am.

See also Reviewer 1 point 3. And the general comment below (Point 8). Practically, on average, if our goal is to predict 11 days ahead, when the AR model is at its best, CHNG-CLI doesn't help (at least with these simple models). However, CHNG-COVID, DV-CLI, and Google-AA “buy” 4 extra days for the same accuracy. CTIS-CLIIC “buys” about 5 extra days. What we mean here is that if we have a certain tolerance for inaccuracy, adding these signals increases the horizon at which we can predict within our tolerance. As for the practical significance of 4-5 extra days, this would of course depend on the use case. However, cases grow exponentially, so making public policy changes 15 or 16 days ahead rather than 11 can mean real gains. There also seems to be a fairly constant mapping between cases and future hospitalizations. Giving hospitals a few extra days to accumulate supplies and find personnel is meaningful. For predicting hotspots, we see similar results (gains of a few days) though with different indicators.

4. For the cumsum plots I mentioned above, and other related techniques, we can plot all the model prediction error cumsum curves versus time and thereby compare different models for size, nonstationarity and significance of prediction improvement. The plots that have been presented in Figure 4 say do none of those tasks.

TODO Any ideas here? It's not clear to me that this is different from above.

5. The idea of using Asof data is of course important but I have known about it for twenty years and practiced it routinely; it has apparently been in routine use in econometrics for 40 years.

We appreciate the reviewer's perspective on the ubiquity of Asof data in econometrics. But the epidemic modelling community is less aware. And it has a real and meaningful importance for model selection going forward. We hope that this paper will lead to better behaviour in this community.

Comments on Significance Statement:

6. I don't know what "all provide a nontrivial boost in accuracy", really means. I guess that the general audience won't either.

We have changed the clause to "all provide a nontrivial improvement in forecast accuracy"

7. The sentence following is missing a word: depends ***on*** the pandemic...

Fixed. Thank you for catching this.

General Comments:

8. In my understanding there are some standard models which have been relied upon I guess millions of times during the pandemic. For example IHME's model. The authors say a sentence or so about their choice of baseline, saying what seems to me quite vague about the relation between their baseline and the heavily used existing models. Wouldn't it be better to identify a specific widely-used model explicitly as a baseline and then develop measures of improvement over that specific widely used model? Also, might the authors explain how the improvements observed by the authors stack up in terms of concrete differences over the widely used models that people have been using. Are there major misses that have taken place and could have been avoided?

We have tried to focus mainly on simple time series models in this paper rather than comparing to other more elaborate models. A careful (out-of-sample) comparison of forecasting performance for state-of-the-art models can be found in reference 6 (Cramer et al.). That paper focuses on forecasting deaths rather than cases and at the state-level rather than disaggregated HRRs. However, their ranking (Figure 2) puts the “COVIDhub-Baseline” about 10th out of 27 models. As we discuss in the manuscript, that baseline is the same as the baseline we use for forecasting in this paper: it propagates today forward and resamples the residuals to create quantiles (we have expanded the description on page 6 to be more explicit). All our “simplistic” models beat this baseline (the AR(3) does as well). For comparison, IHME (along with 17 other “state-of-the-art” models) does not. We would argue that using any of the models in this paper would be preferable to relying on many of the “widely-used” models if the goal is forecast accuracy.

- We have added a sentence to the introduction (around line 40) that describes the performance of simple models relative to widely-used models as shown in Cramer et al.
- We have expanded our description of our baseline model (around line 457). In particular, we emphasize that it is the same as that used by the COVID-19 Forecast Hub.